Post-test exploratory abstention: no threshold selected

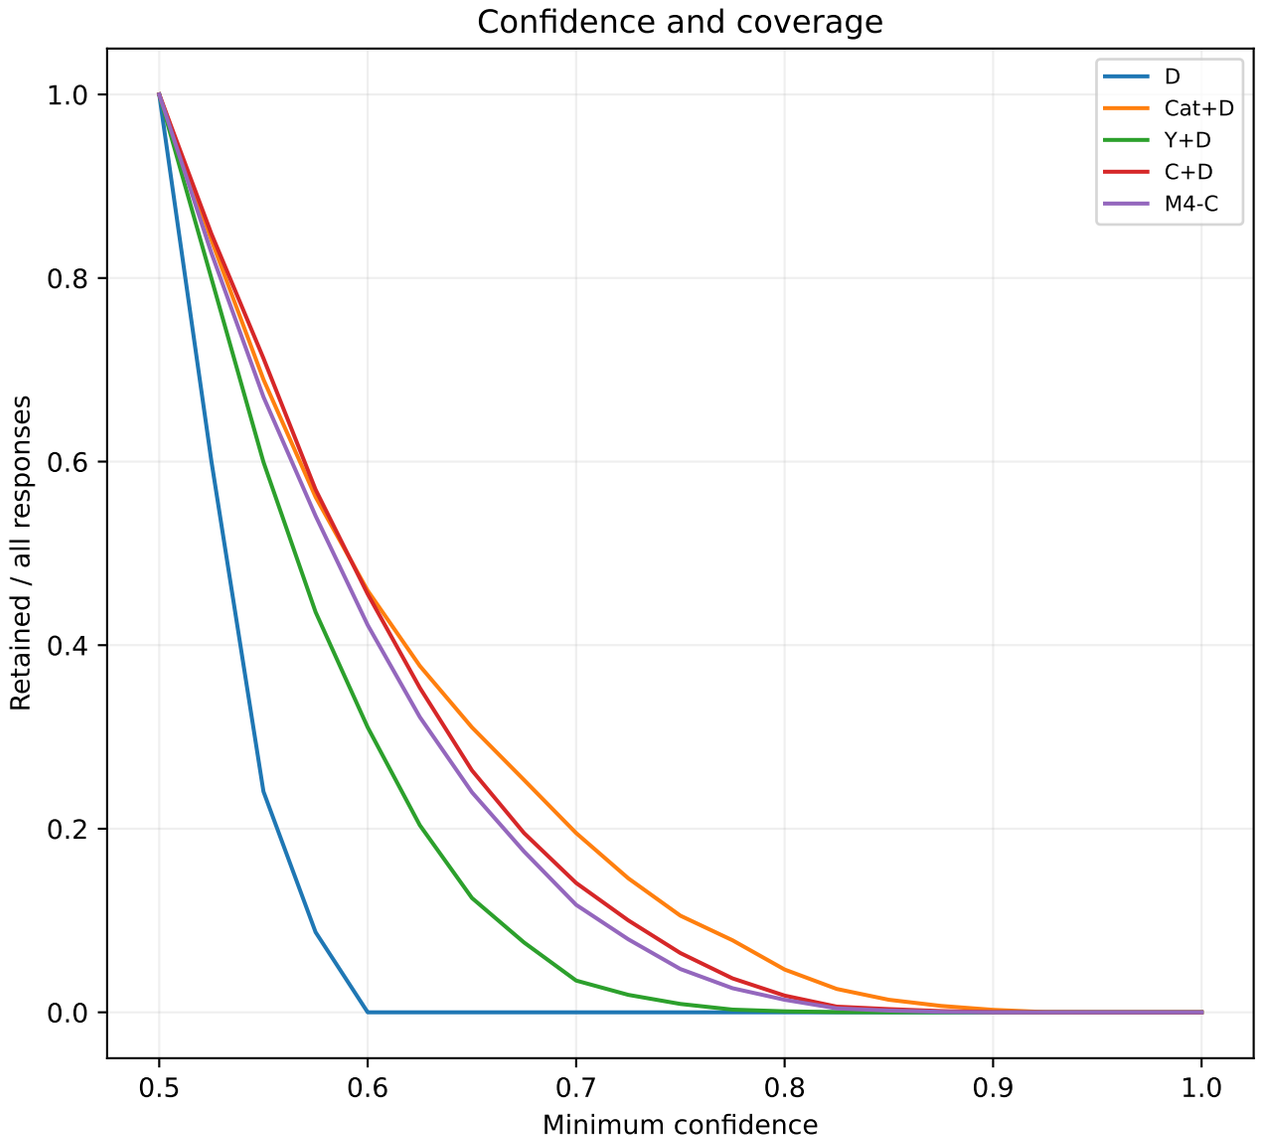

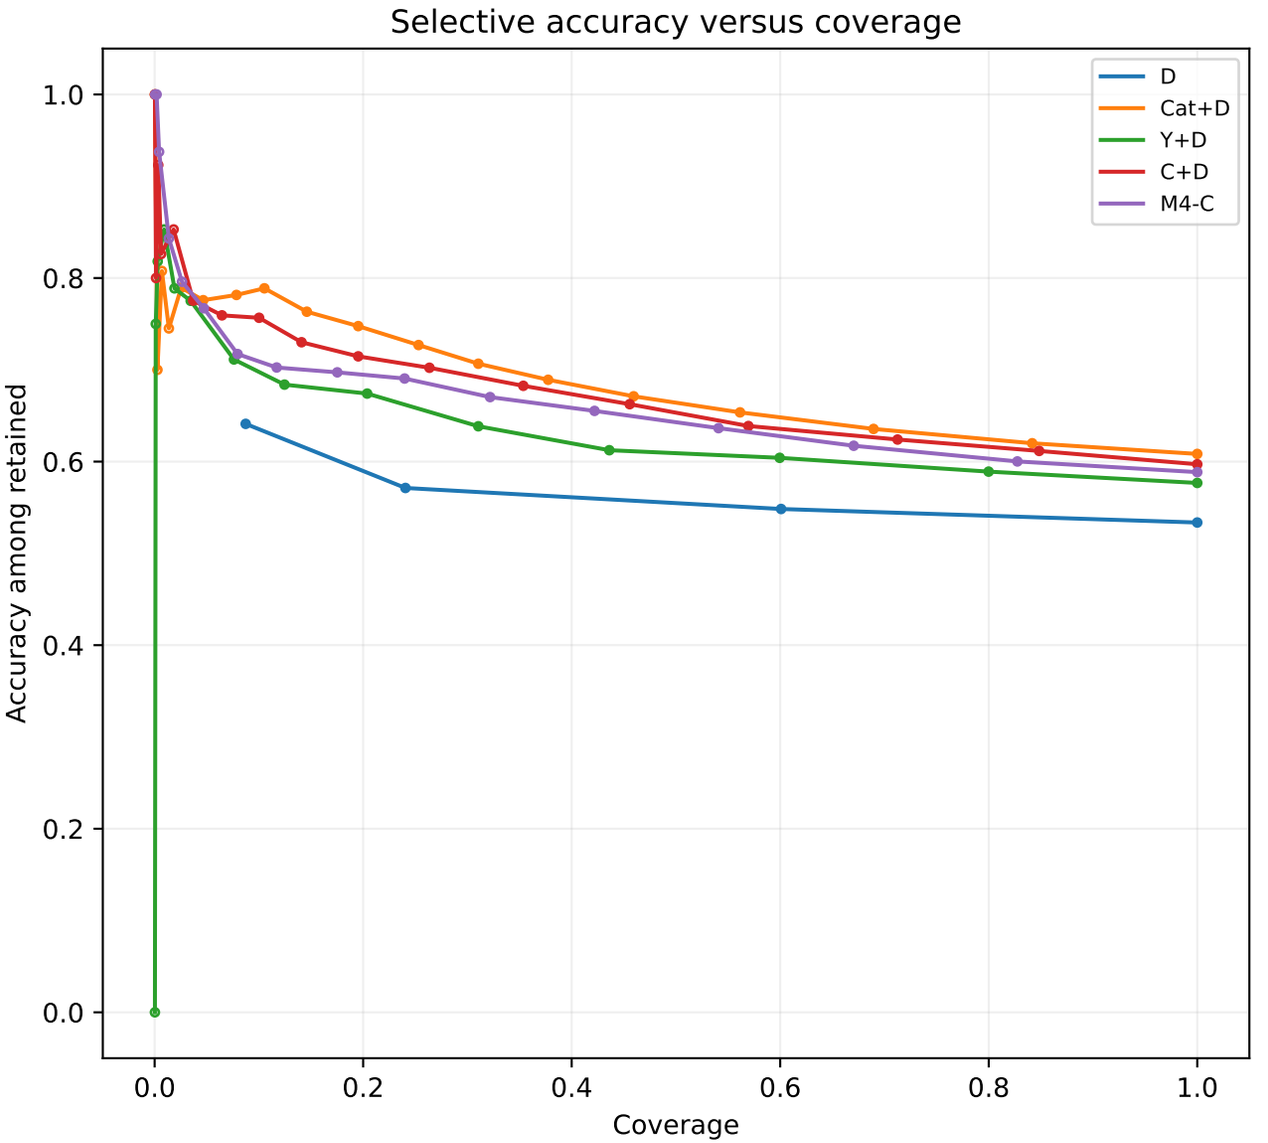

Hollow points: <100 retained responses or <30 pairs. No retained cases: accuracy undefined.

Utility includes abstentions in the denominator

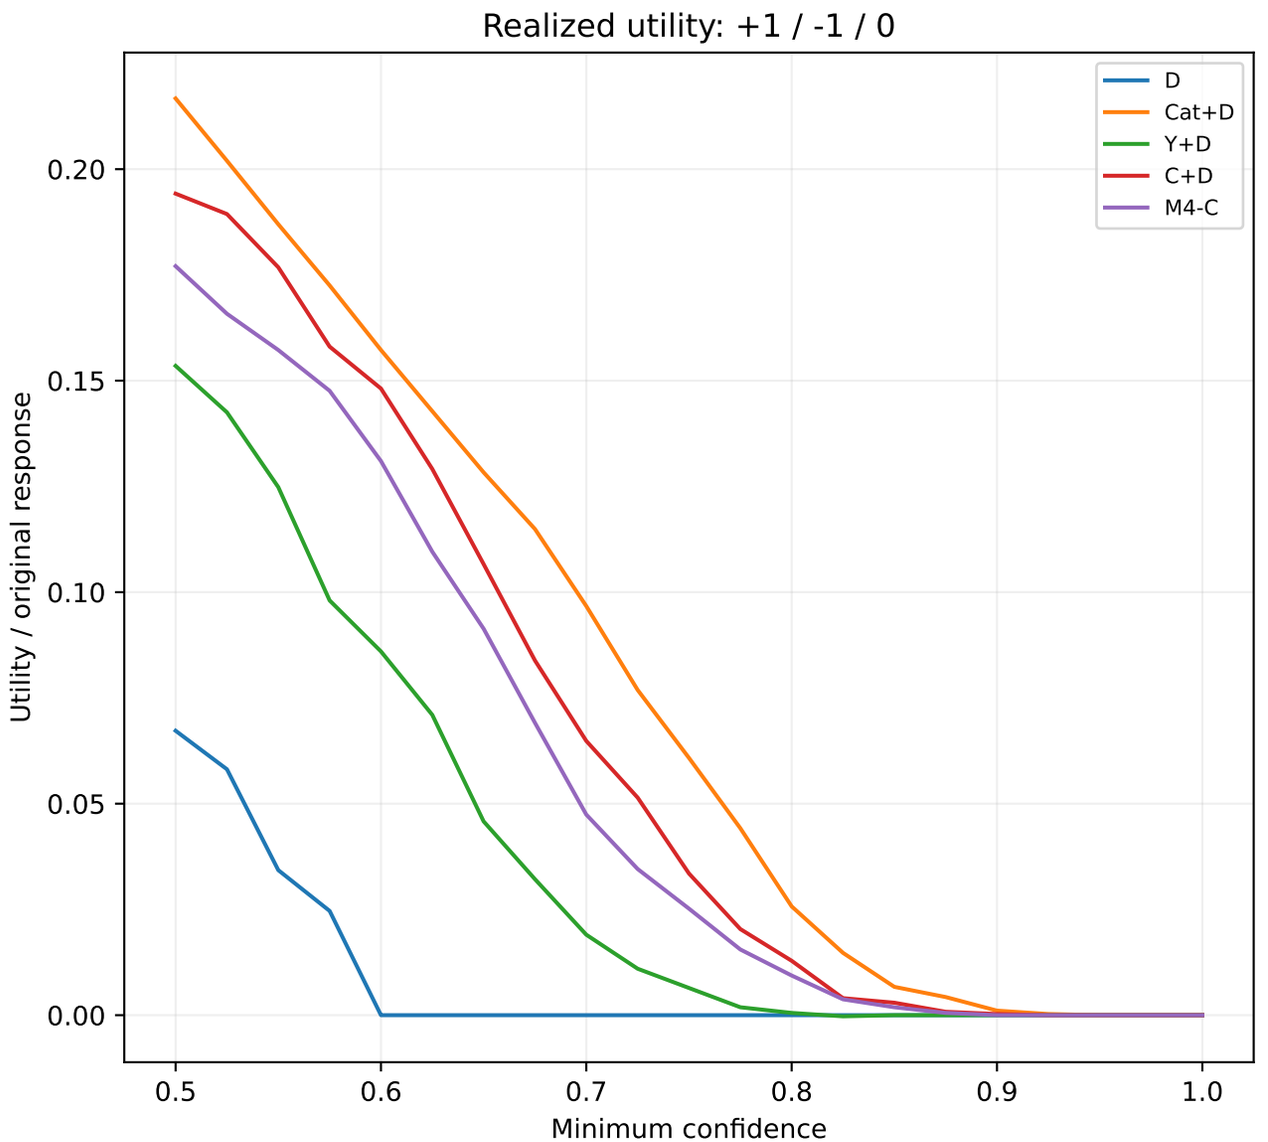

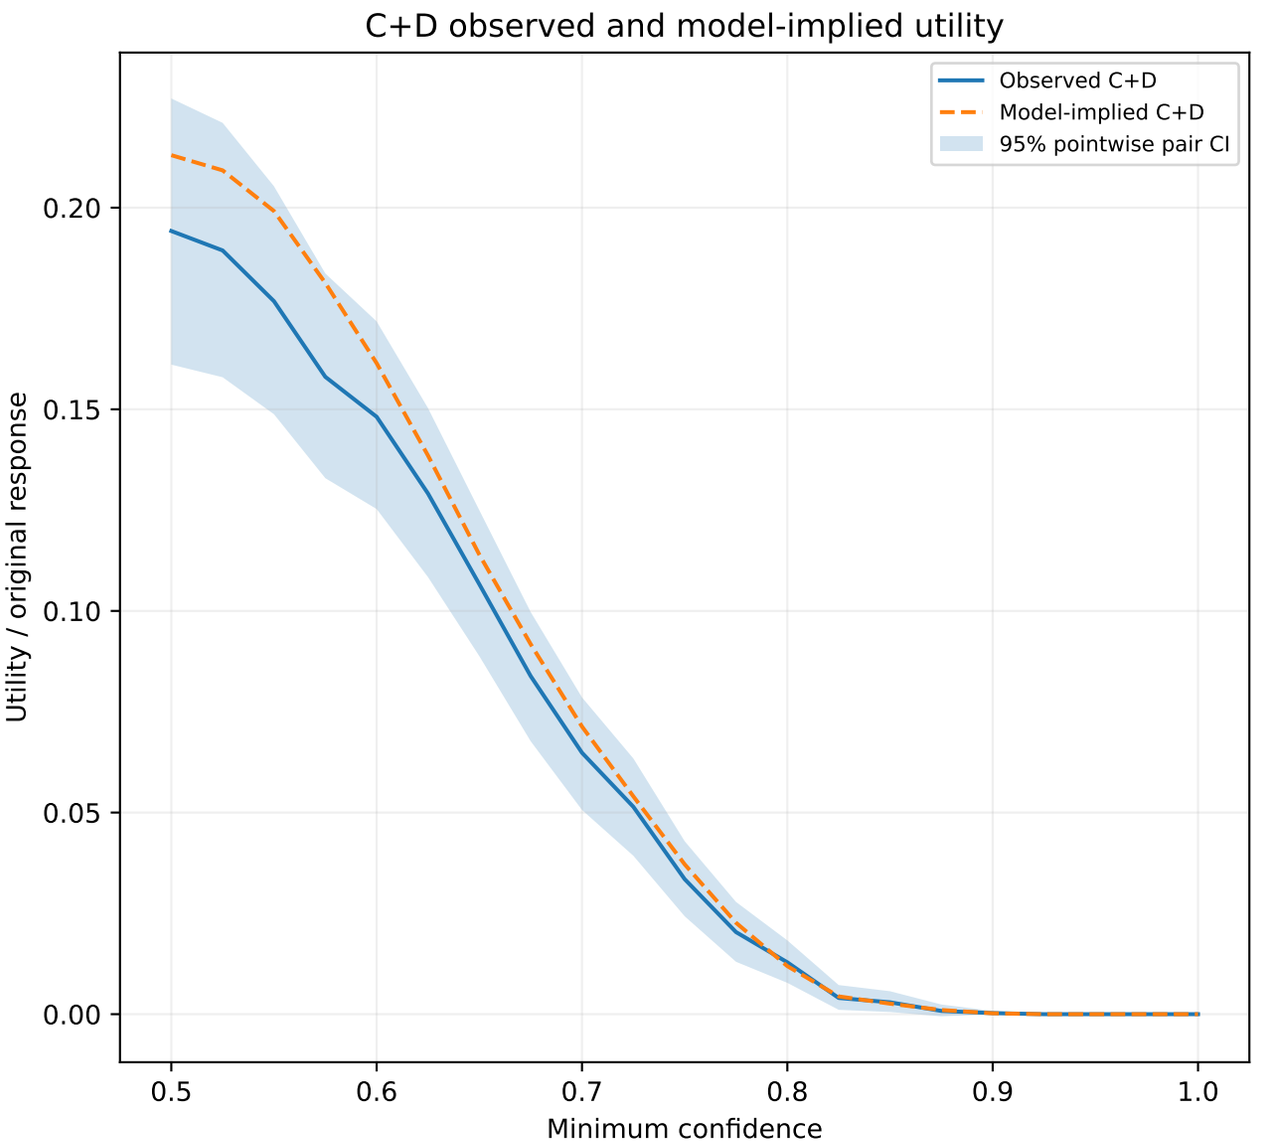

Implied utility is expected reward only if probabilities are calibrated; intervals are not multiplicity adjusted.

Fixed illustrative cutoffs: counts and support

Model	Cutoff	Retained	Pairs	Coverage	Accuracy	Utility / all	Support
D	0.50	3733	1207	100.0%	53.4%	0.0672	supported
D	0.60	0	0	0.0%	undefined	0.0000	no_retained_cases
D	0.70	0	0	0.0%	undefined	0.0000	no_retained_cases
Cat+D	0.50	3733	1207	100.0%	60.8%	0.2167	supported
Cat+D	0.60	1715	717	45.9%	67.1%	0.1572	supported
Cat+D	0.70	729	306	19.5%	74.8%	0.0967	supported
Y+D	0.50	3733	1207	100.0%	57.7%	0.1535	supported
Y+D	0.60	1159	627	31.0%	63.8%	0.0860	supported
Y+D	0.70	129	95	3.5%	77.5%	0.0190	supported
C+D	0.50	3733	1207	100.0%	59.7%	0.1942	supported
C+D	0.60	1701	768	45.6%	66.3%	0.1481	supported
C+D	0.70	526	260	14.1%	73.0%	0.0648	supported
M4-C	0.50	3733	1207	100.0%	58.9%	0.1771	supported
M4-C	0.60	1575	619	42.2%	65.5%	0.1310	supported
M4-C	0.70	437	180	11.7%	70.3%	0.0474	supported

These cutoffs were fixed for description, not chosen to maximize test utility.

Interpretation and safeguards

- Post-test retrospective analysis; the saved G2 outcomes were already exposed.

- No model or deployment threshold is selected from these curves.

- Correct +1, incorrect -1, abstain 0; utility denominator is all original responses.

- Retained accuracy is undefined when no predictions remain.

- 5,000 shared unordered-pair bootstrap draws; 95% pointwise intervals are unadjusted.

- Respondent linkage, shared-image dependence and model-selection uncertainty remain unaccounted for.

- Under calibrated probabilities and these payoffs, always predict is weakly optimal in expected total utility.

- Demographics-only predicts displayed A; A is not established to mean left.