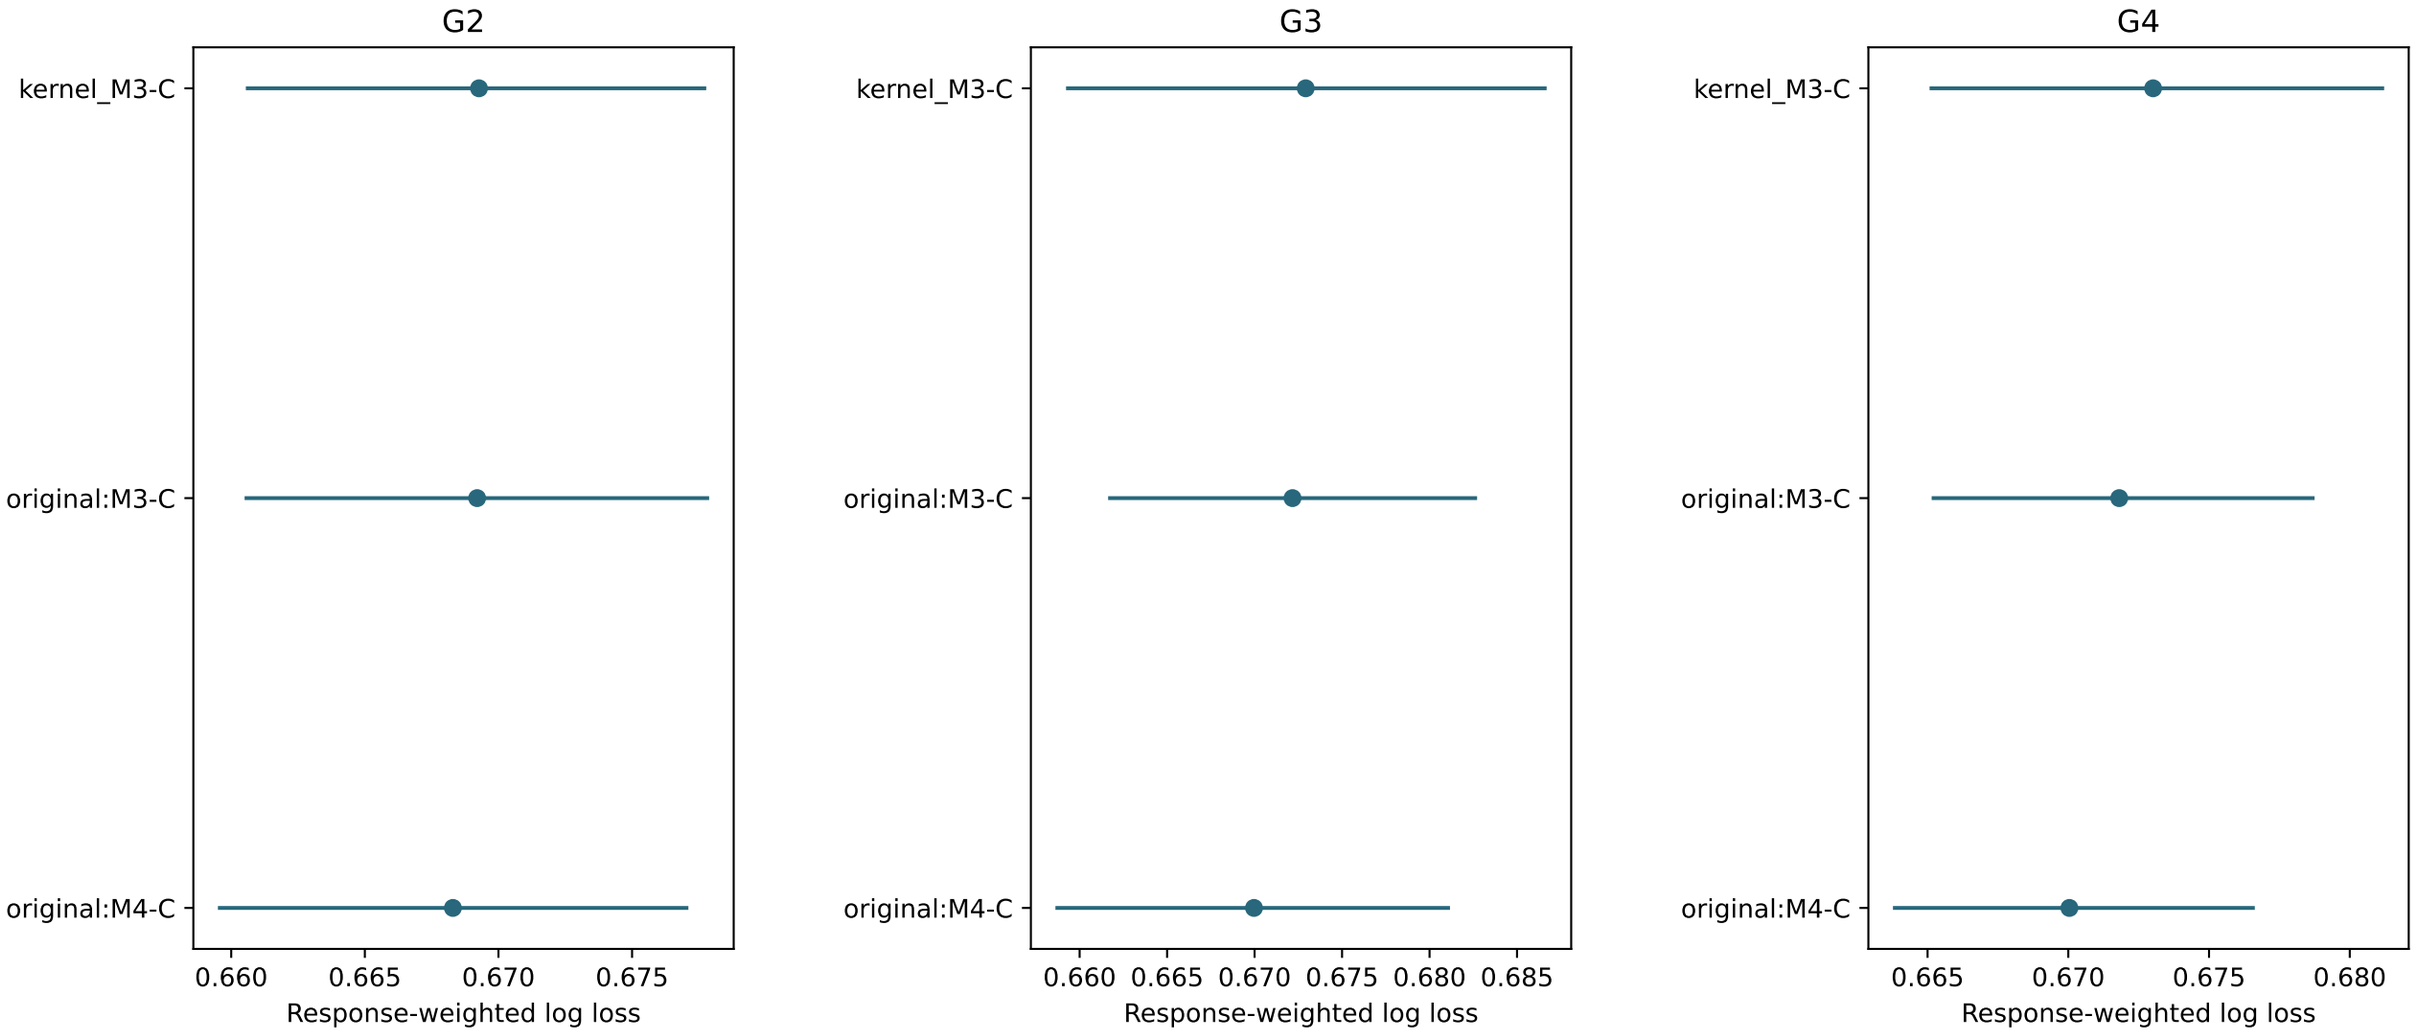
Pooled nonlinear CNN versus linear scorers
G2
G3
G4
kernel_M3-C
original:M3-C
original:M4-C
0.660
0.665
0.670
0.675
0.680
0.685
Response-weighted log loss
Retrospective, post-test, not causal or independent; 95% pointwise pair/photo/cat bootstrap.

Fit-only transformations; 1,836 common validation responses

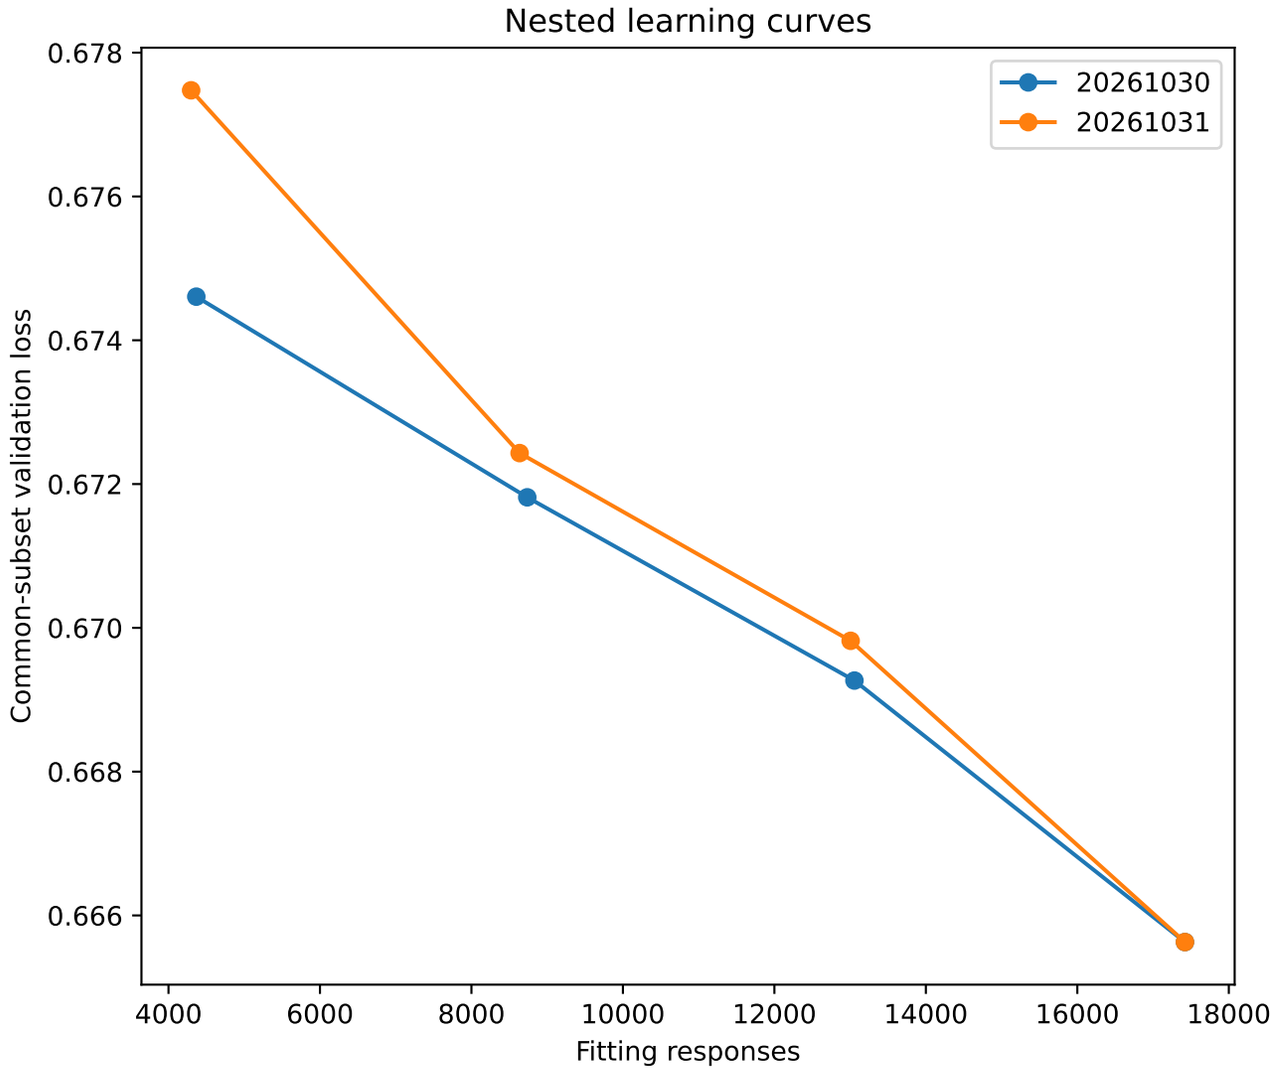

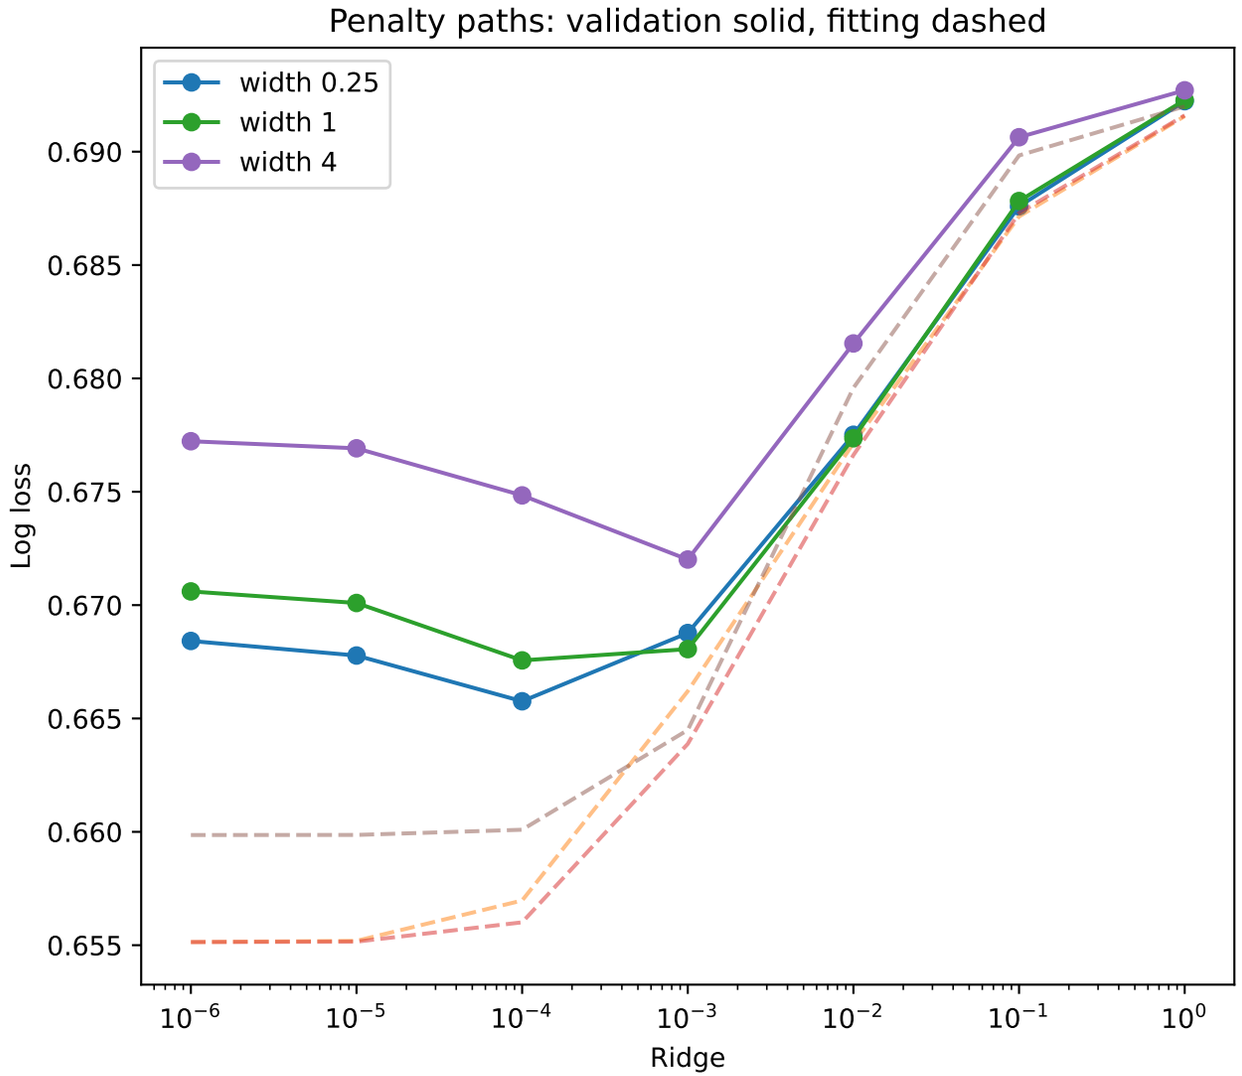

All 21 width/ridge candidates refit per subset; the two 100% seed inputs coincide.

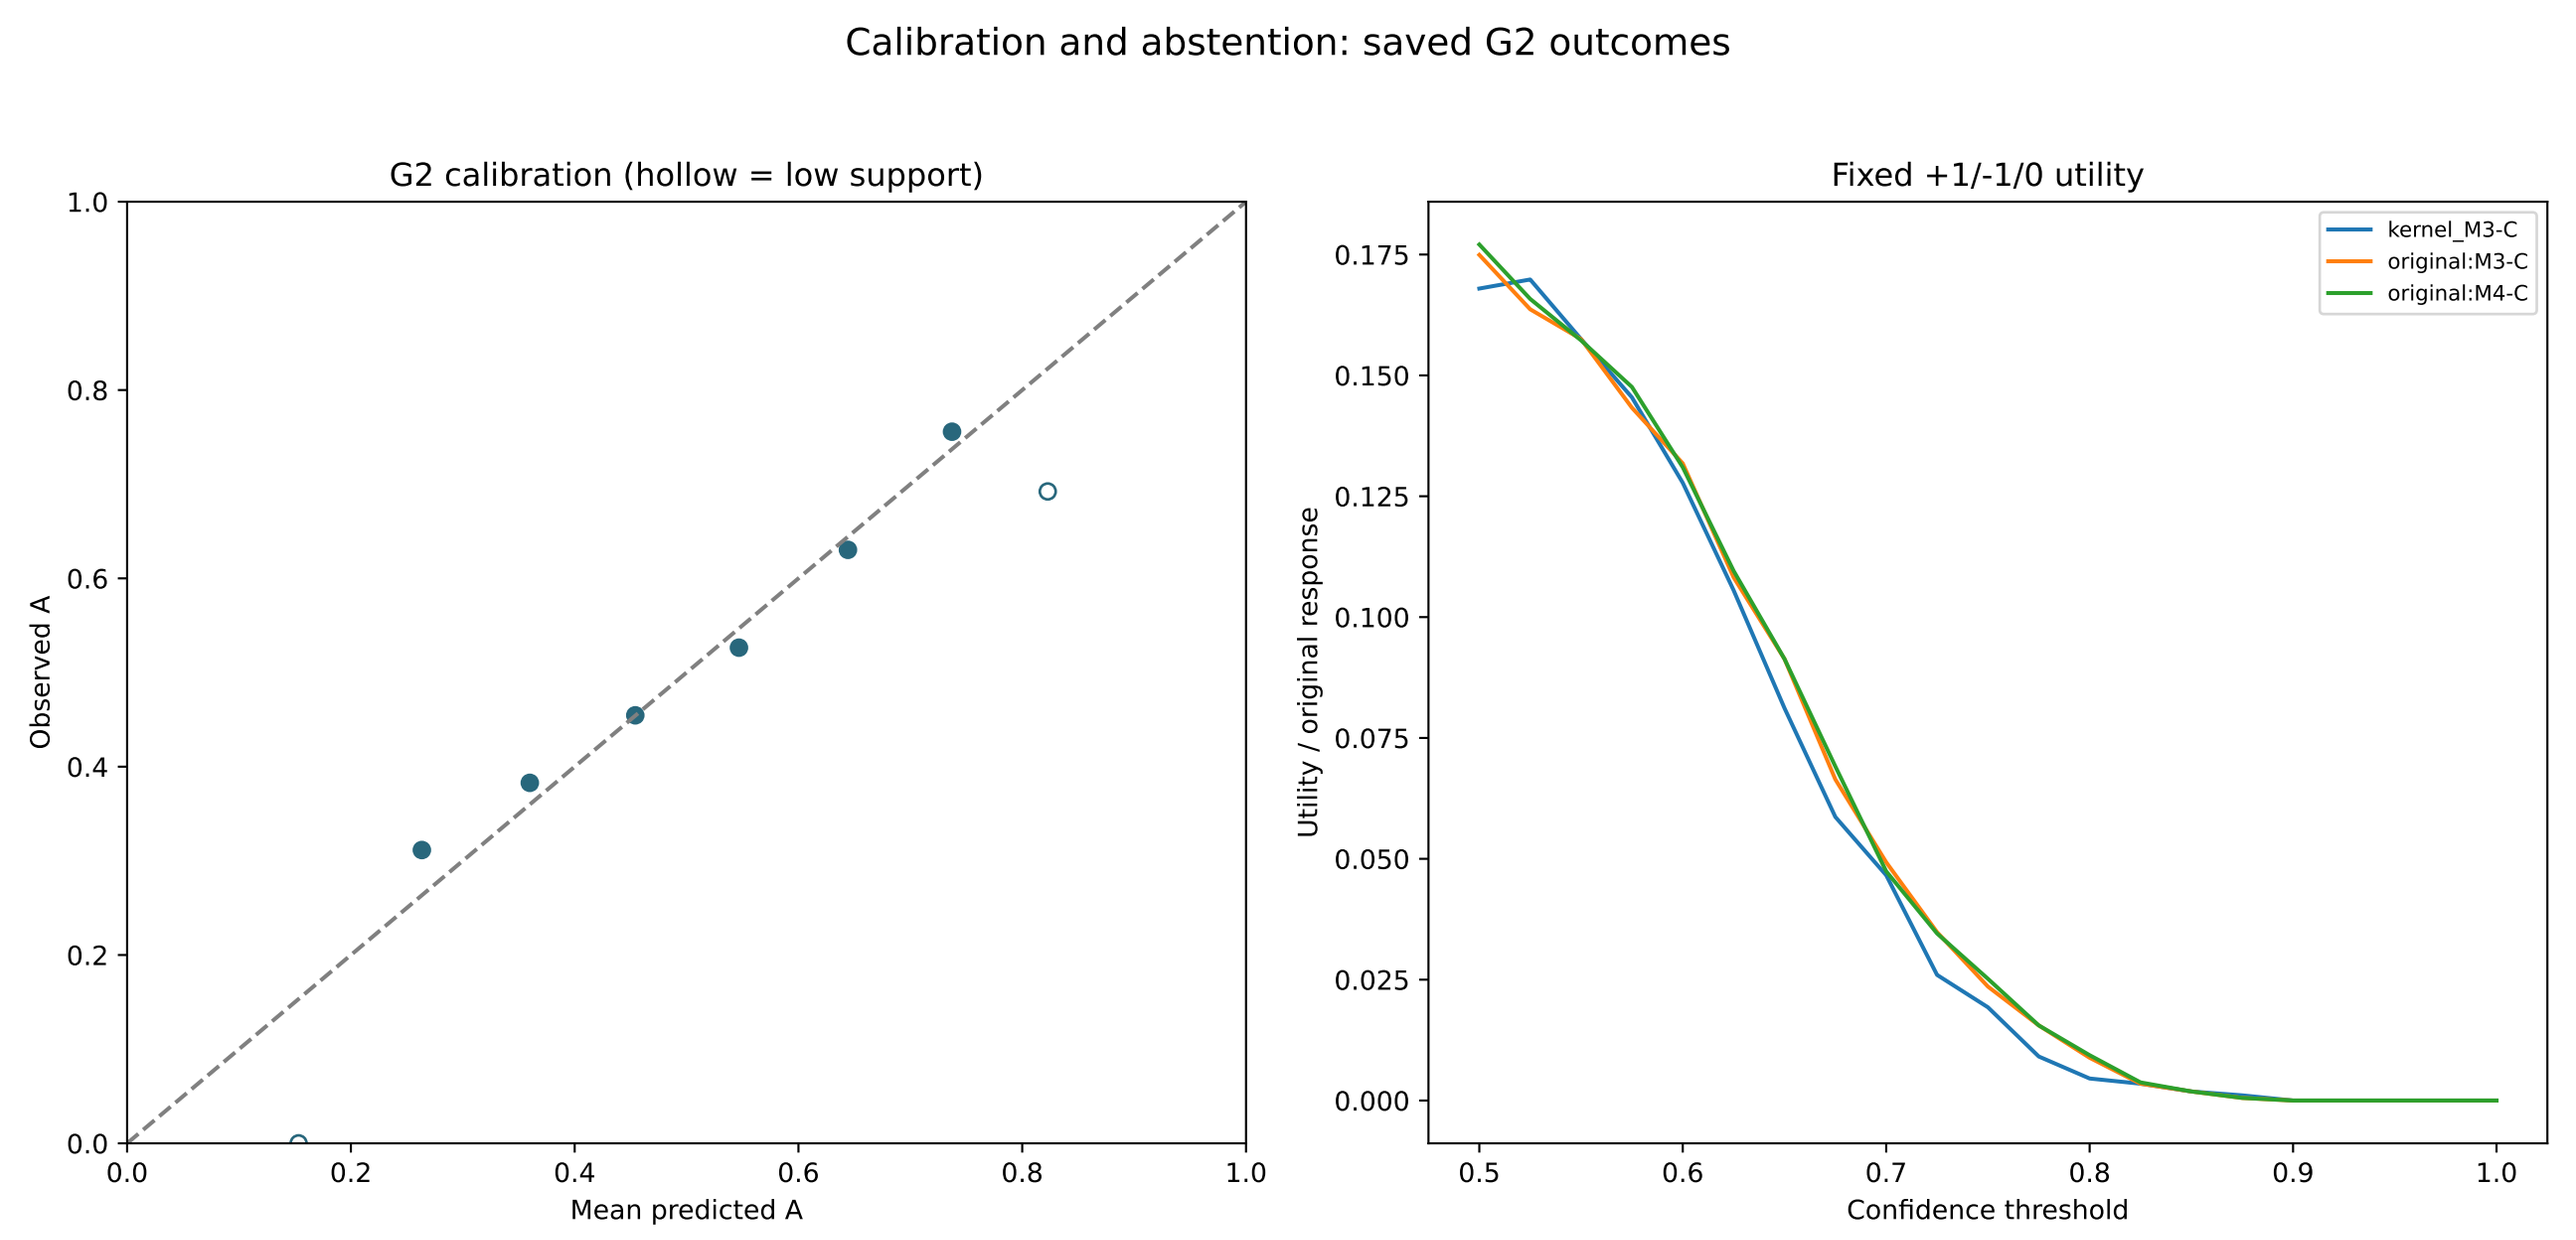

Post-test descriptive only; no threshold or deployment model selected.

Interpretation and safeguards

- Predictive association, not causal effect.

- Original G2, G3 and G4 outcomes were already exposed: no independent confirmation.

- The scorer is one nonlinear scalar per photo, not an unrestricted pair model.

- G3 and G4 use separately refitted fitting-only landmarks and validation-tuned settings.

- Pointwise 95% intervals omit model-selection and some shared-image/respondent dependence.

- No post-test model repair, official-prediction change or abstention threshold selection.